


# Trustworthy AI Systems

-- Image Classification

Instructor: Guangjing Wang  
[guangjingwang@usf.edu](mailto:guangjingwang@usf.edu)


# Trustworthy AI Systems



Application areas:

- Vision Domain
- Audio Domain
- Text Domain

# Classical Computer Vision Tasks




Localizing and labeling objects

Dividing images into regions

# Traditional Computer Vision: Feature Engineering

- Feature Engineering: preprocesses raw data by transforming and selecting relevant features.

- Histogram: the occurrences of each pixel intensity value. This is an important feature for object recognition.



Example: Histogram

# Data-driven Computer Vision

1. Collect a set of images and labels
2. Use deep learning algorithms to train a classifier or regression model
3. Evaluate the model on new images

```
def train(images, labels):
    # Machine learning!
    return model
```

```
def predict(model, test_images):
    # Use model to predict labels
    return test_labels
```

**airplane**



**automobile**



**bird**




**cat**



**deer**



# Data-driven Computer Vision




# Deep Learning for Image Classification




<https://stock.adobe.com/search?k=panda>


```
1 import torch
2 import torch.nn as nn
3 class Model4_1(nn.Module):
4     def __init__(self):
5         super(Model4_1, self).__init__()
6         self.lin1 = nn.Linear(784, 100)
7         self.relu = nn.ReLU()
8         self.lin2 = nn.Linear(100, 10)
9
10    def forward(self, x):
11        out = self.lin1(x)
12        out = self.relu(out)
13        out = self.lin2(out)
14        return out
15
16 model4_1 = Model4_1()
```




# Deep Learning: A general term for various DNNs



<https://becominghuman.ai/building-a-convolutional-neural-network-cnn-model-for-image-classification-116f77a7a236>




<https://machinelearningmastery.com/the-transformer-model/>




<https://www.analyticsvidhya.com/blog/2022/03/a-brief-overview-of-recurrent-neural-networks-rnn/>



# Why Deep Learning?



# Convolutional Layer




activation map



<https://hannibunny.github.io/mlbook/neuralnetworks/convolutionDemos.html>

$$n_{out} = \left\lfloor \frac{n_{in} + 2p - k}{s} \right\rfloor + 1$$

# Convolutional Neural Network



# Conv Layer in PyTorch

## Conv2d

```
CLASS torch.nn.Conv2d(in_channels, out_channels, kernel_size, stride=1, padding=0, dilation=1,
groups=1, bias=True, padding_mode='zeros', device=None, dtype=None) [SOURCE]
```

Applies a 2D convolution over an input signal composed of several input planes.

In the simplest case, the output value of the layer with input size  $(N, C_{\text{in}}, H, W)$  and output  $(N, C_{\text{out}}, H_{\text{out}}, W_{\text{out}})$  can be precisely described as:

$$\text{out}(N_i, C_{\text{out},j}) = \text{bias}(C_{\text{out},j}) + \sum_{k=0}^{C_{\text{in}}-1} \text{weight}(C_{\text{out},j}, k) * \text{input}(N_i, k)$$


where  $*$  is the valid 2D [cross-correlation](#) operator,  $N$  is a batch size,  $C$  denotes a number of channels,  $H$  is a height of input planes in pixels, and  $W$  is width in pixels.

This module supports [TensorFloat32](#).

On certain ROCm devices, when using float16 inputs this module will use [different precision](#) for backward.

- `stride` controls the stride for the cross-correlation, a single number or a tuple.
- `padding` controls the amount of padding applied to the input. It can be either a string {'valid', 'same'} or an int / a tuple of ints giving the amount of implicit padding applied on both sides.
- `dilation` controls the spacing between the kernel points; also known as the u00e0 trous algorithm. It is harder to describe, but this [link](#) has a nice visualization of what `dilation` does.


# ConvNet JS Demo




<https://cs.stanford.edu/people/karpathy/convnetjs/demo/cifar10.html>

# Summary: components of CNNs


Convolution Layers




Pooling Layers




Fully-Connected Layers



DNN Example




Activation Function



Normalization

$$\hat{x}_{i,j} = \frac{x_{i,j} - \mu_j}{\sqrt{\sigma_j^2 + \epsilon}}$$

Max Pooling Example



# Batch Normalization: normalizing inputs to each layer

Consider a single layer  $y = Wx$


The following could lead to tough optimization:

- Inputs  $x$  are not *centered around zero* (need large bias)
- Inputs  $x$  have different scaling per-element  
(entries in  $W$  will need to vary a lot)

Idea: force inputs to be “nicely scaled” at each layer!

# Batch Normalization: normalizing inputs to each layer

**Input:**  $x : N \times D$



$$\mu_j = \frac{1}{N} \sum_{i=1}^N x_{i,j} \quad \text{Per-channel mean, shape is } D$$

$$\sigma_j^2 = \frac{1}{N} \sum_{i=1}^N (x_{i,j} - \mu_j)^2 \quad \text{Per-channel var, shape is } D$$

$$\hat{x}_{i,j} = \frac{x_{i,j} - \mu_j}{\sqrt{\sigma_j^2 + \varepsilon}} \quad \text{Normalized x, Shape is } N \times D$$

Problem: What if zero-mean, unit variance is too hard of a constraint?

# Batch Normalization: normalizing inputs to each layer

**Input:**  $x : N \times D$

**Learnable scale and shift parameters:**

$\gamma, \beta : D$

Learning  $\gamma = \sigma$ ,  
 $\beta = \mu$  will recover the identity function!

$$\mu_j = \frac{1}{N} \sum_{i=1}^N x_{i,j} \quad \text{Per-channel mean, shape is D}$$

$$\sigma_j^2 = \frac{1}{N} \sum_{i=1}^N (x_{i,j} - \mu_j)^2 \quad \text{Per-channel var, shape is D}$$

$$\hat{x}_{i,j} = \frac{x_{i,j} - \mu_j}{\sqrt{\sigma_j^2 + \varepsilon}} \quad \text{Normalized x, Shape is } N \times D$$

$$y_{i,j} = \gamma_j \hat{x}_{i,j} + \beta_j \quad \text{Output, Shape is } N \times D$$

# Batch Normalization: Test Time

**Input:**  $x : N \times D$

$\mu_j =$  (Running) average of values seen during training

Per-channel mean, shape is D

**Learnable scale and shift parameters:**

$\gamma, \beta : D$

$\sigma_j^2 =$  (Running) average of values seen during training

Per-channel var, shape is D

During testing batchnorm becomes a linear operator!  
Can be fused with the previous fully-connected or conv layer

$$\hat{x}_{i,j} = \frac{x_{i,j} - \mu_j}{\sqrt{\sigma_j^2 + \varepsilon}}$$


Normalized x,  
Shape is N x D

$$y_{i,j} = \gamma_j \hat{x}_{i,j} + \beta_j$$





Output,  
Shape is N x D

# Not homework... but read papers to learn

- Why using normalization?
- Other normalization techniques?
- N: batch size, C: channel size, H,W: height and width




# CNN Architectures




# ResNet (1)

Very deep networks using residual connections:

- 152-layer model for ImageNet
- ILSVRC'15 classification winner (3.57% top 5 error) - Swept all classification and detection competitions in ILSVRC'15 and COCO'15!




# ResNet (2)




# ResNet (3)

- Total depths of 18, 34, 50, 101, or 152 layers for ImageNet
  - Stack residual blocks
  - Every residual block has two 3x3 conv layers
  - Periodically, double number of filters and down sample spatially using stride 2 (/2 in each dimension)
  - Additional conv layer at the beginning
  - No FC layers at the end (only FC 1000 to output classes)





# ResNet (4)

For deeper networks (ResNet-50+), use “bottleneck” layer to improve efficiency (similar to GoogLeNet)



# Why ResNet



Problem: Deeper models are harder to optimize

Solution: Copying the learned layers from the shallower model and setting additional layers to identity mapping

# ResNet Practice


Training ResNet in practice:

- Batch Normalization after every CONV layer
- Xavier initialization from He et al.
- Stochastic Gradient Descent (SGD) + Momentum (0.9)
- Learning rate: 0.1, divided by 10 when validation error plateaus
- Mini-batch size 256
- Weight decay of 1e-5
- No dropout used


# Back to Classification: Challenges

## Variations in the physical world

- Illumination
- Background Clutter
- Occlusion
- Deformation
- Intraclass variation




# Challenges in Classification: Context



[https://www.linkedin.com/posts/ralph-aboujaoude-diaz-40838313\\_technology-artificialintelligence-computervision-activity-6912446088364875776-h-lq/?utm\\_source=linkedin\\_share&utm\\_medium=member\\_desktop\\_web](https://www.linkedin.com/posts/ralph-aboujaoude-diaz-40838313_technology-artificialintelligence-computervision-activity-6912446088364875776-h-lq/?utm_source=linkedin_share&utm_medium=member_desktop_web)

# For your project: Data Preprocessing



# For your project: Transfer Learning


Have some dataset of interest but it has < ~1M images?

- Find a very large dataset that has similar data, train a big model there
- Transfer the learning to your dataset

Deep learning frameworks provide a “Model Zoo” of pretrained models so you don’t need to train your own

- <https://github.com/tensorflow/models>
- <https://github.com/pytorch/vision>


# For your project: Transfer Learning



	<b>very similar dataset</b>	<b>very different dataset</b>
<b>very little data</b>	Use Linear Classifier on top layer	You're in trouble... Try linear classifier from different stages
<b>quite a lot of data</b>	Finetune a few layers	Finetune a larger number of layers or start from scratch!


# For your project: Transfer Learning

Object Detection  
(Fast R-CNN)



CNN pretrained  
on ImageNet

Image Captioning: CNN + RNN




# For your project: Some Practices

Consider CIFAR-10 example with [32,32,3] images:

- Data Preprocessing:
  - Subtract the mean image (e.g. AlexNet) (mean image = [32,32,3] array)
  - Subtract per-channel mean (e.g. VGGNet) (mean along each channel = 3 numbers)
  - Subtract per-channel mean and Divide by per-channel std (e.g. ResNet and beyond) (mean along each channel = 3 numbers)
- Weight Initialization: Kaiming / MSRA Initialization
- Use ReLU. Be careful with your learning rates
- Try out Leaky ReLU / PReLU / GELU (**Check them out by yourself**)


# For your project: Early Stopping



Stop training the model when accuracy on the validation set decreases. Train for a long time, but always keep track of the model snapshot that worked best on val.

# For your project: Model Ensembles

- Train multiple independent models
- At test time average their results



<https://pub.towardsai.net/introduction-to-ensemble-methods-226a5a421687>

# For your project: Regularization (1)

- Add a term to a loss:

$$L = \frac{1}{N} \sum_{i=1}^N \sum_{j \neq y_i} \max(0, f(x_i; W)_j - f(x_i; W)_{y_i} + 1) + \boxed{\lambda R(W)}$$

In common use:

L2 regularization

$$R(W) = \sum_k \sum_l W_{k,l}^2 \quad (\text{Weight decay})$$

L1 regularization

$$R(W) = \sum_k \sum_l |W_{k,l}|$$


Elastic net (L1 + L2)

$$R(W) = \sum_k \sum_l \beta W_{k,l}^2 + |W_{k,l}|$$

- Random Dropout, 0.5 is common


# For your project: Regularization (2)

- Data Augmentation
  - Horizontal Flips
  - Random crops and scales
  - Color Jitter
  - Rotation
  - Shearing
  - ....






# For your project: Regularization (3)

- Training: the core idea is to add random noise
  - Dropout: Consider dropout for large fully connected layers
  - Batch Normalization
  - Data Augmentation
  - Cutout / Random Crop : Try cutout especially for small classification datasets
- Testing: Marginalize over the noise



# For your project: Look at the Learning Curve



# Reference: Stanford Spring 2024 cs231n

- <https://cs231n.stanford.edu/schedule.html>
- [https://cs231n.stanford.edu/slides/2024/lecture\\_5.pdf](https://cs231n.stanford.edu/slides/2024/lecture_5.pdf)
- [https://cs231n.stanford.edu/slides/2024/lecture\\_6\\_part\\_1.pdf](https://cs231n.stanford.edu/slides/2024/lecture_6_part_1.pdf)
- [https://cs231n.stanford.edu/slides/2024/lecture\\_6\\_part\\_2.pdf](https://cs231n.stanford.edu/slides/2024/lecture_6_part_2.pdf)