

# Trustworthy AI Systems

-- Generative Modeling (Part I)

Instructor: Guangjing Wang

[guangjingwang@usf.edu](mailto:guangjingwang@usf.edu)

# Group Member Checkpoint

- For the course projects, two to three students will form a group, choose an AI application topic/task, and complete both midterm and final projects.
- The deadline is Tonight: Sep. 8<sup>th</sup>, 11:59 pm
- If you cannot find your teammates, we will help randomly assign a group, but you will get a 0 grade for this checkpoint.

# Project Examples from Last Semester

- Research-oriented
  - Voice conversion and reversing the converted voice
  - Machine unlearning
  - Enhancing low-light image processing with Retinex-based algorithm
- Engineering-oriented
  - Attendance Record Based on Face Identification
  - LLM Agent for Food Classification
  - AI-assisted symptom analysis for healthcare diagnostics
  - Signature forgery detection using deep learning models

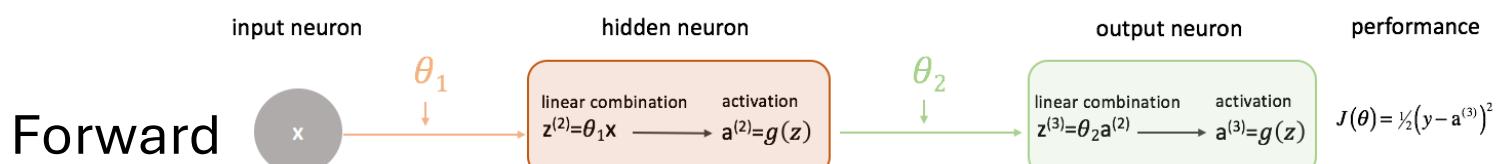
# Last Lecture

- Semantic Segmentation
- Object Detection
  - R-CNN series
  - YOLO series
- Classification and Regression
  - Model structure: ResNet, Unet, Transformer
  - Optimization goal: loss function
  - Hyperparameters for training

# Neural Network Training

Chain Rule

$$\frac{\partial}{\partial z} p(q(z)) = \frac{\partial p}{\partial q} \frac{\partial q}{\partial z}$$



Backward

$$\frac{\partial J(\theta)}{\partial \theta_2} = \left( \frac{\partial J(\theta)}{\partial a^{(3)}} \right) \left( \frac{\partial a^{(3)}}{\partial z} \right) \left( \frac{\partial z}{\partial \theta_2} \right)$$

```
# forward + backward + optimize
outputs = net(inputs)
loss = criterion(outputs, labels)
loss.backward()
optimizer.step()
```

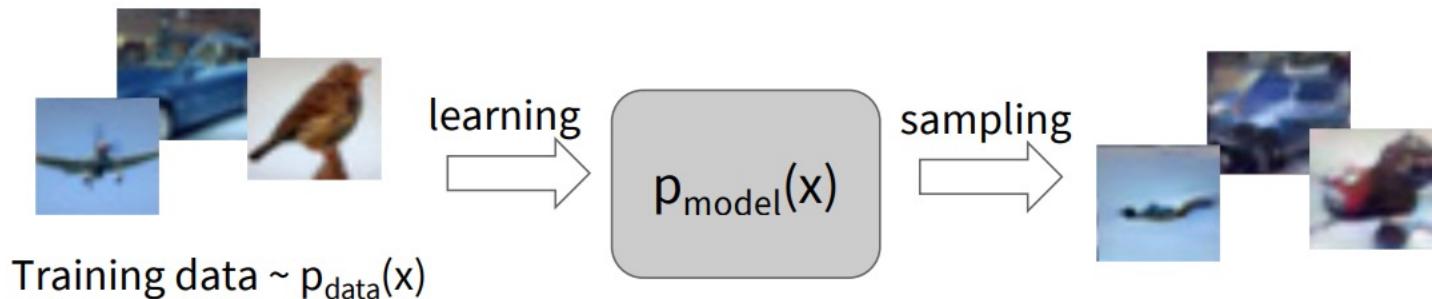
<https://www.jeremyjordan.me/neural-networks-training/>  
[https://docs.pytorch.org/tutorials/beginner/blitz/cifar10\\_tutorial.html](https://docs.pytorch.org/tutorials/beginner/blitz/cifar10_tutorial.html)

# This Lecture

- Generative Modeling
- Generative Adversarial Network (GAN)
  - DCGAN
  - Conditional GAN
  - CycleGAN
- Neural Style Transfer

# Generative Modeling

Given training data, generate new samples from same distribution



Objectives:

1. Learn  $p_{\text{model}}(x)$  that approximates  $p_{\text{data}}(x)$
2. Sampling new  $x$  from  $p_{\text{model}}(x)$

# Learn Data Distributions

- Minimizing certain divergence metrics (e.g., KL divergence) between **the training data distribution**, and **the distribution that the model learns**.
- Training models that maximize the expected log likelihood of  $p_\theta(x)$ 
  - If I sample from the distribution and get a
    - high likelihood → likely the sample came from the training distribution
    - low likelihood → the sample probably did not come from the training distribution

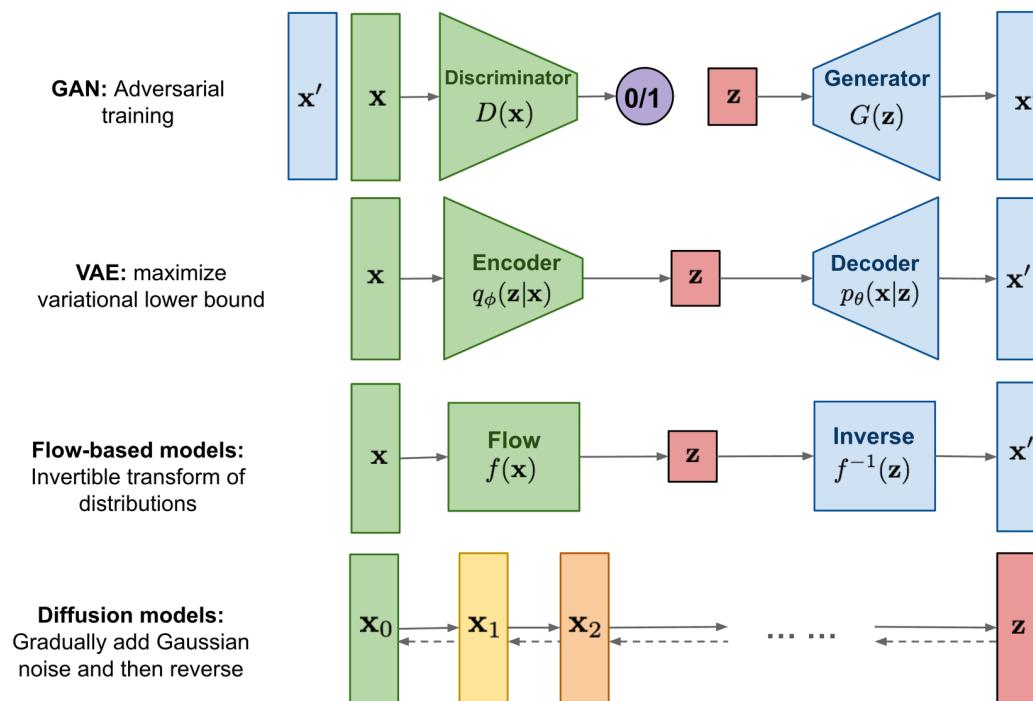
# Why Generative Modeling?

- Realistic samples for artwork, super-resolution, colorization, etc.
- Learn useful features for downstream tasks such as classification.
- Getting insights from high-dimensional data (physics, medical imaging, etc.)
- Modeling physical world for simulation and planning (robotics and reinforcement learning applications)
- Many more ...

# Overview of different types of generative models

GAN/VAE: Implicitly learn probability density function of  $p(\mathbf{x}) = \int p(\mathbf{x}|\mathbf{z})p(\mathbf{z})d\mathbf{z}$

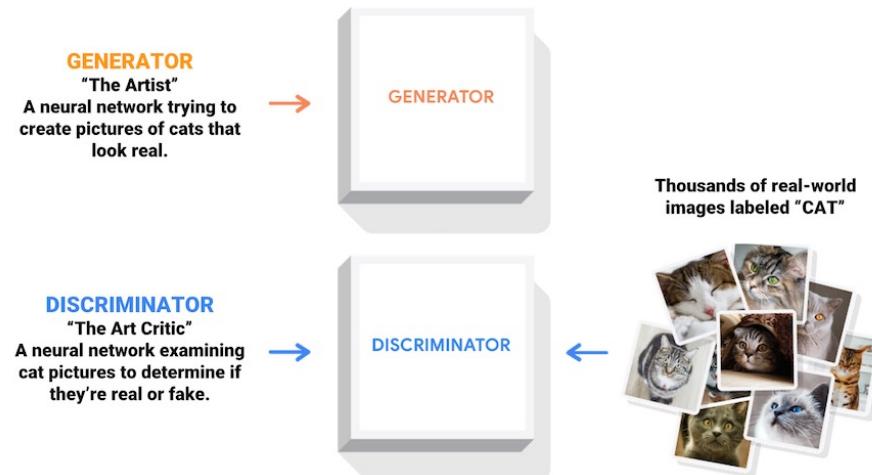
Normalizing-flows: statistics tool for density estimation



Source: <https://lilianweng.github.io/posts/2021-07-11-diffusion-models/>

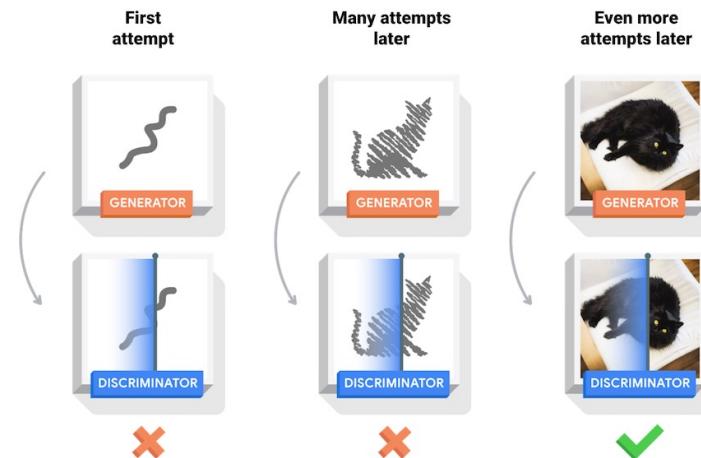
# Generative Adversarial Network (GAN)

- Two models are trained simultaneously by an adversarial process.
  - A generator ("the artist") learns to create images that look real
  - A discriminator ("the art critic") learns to tell real images apart from fakes.



# The idea of GAN

- During training, the generator progressively becomes better at creating images that look real, while the discriminator becomes better at telling them apart.
- The process reaches equilibrium when the *discriminator* can no longer distinguish real images from fakes.



# Deep Convolutional GAN (DCGAN)

- Generator: **Upsampling layers** to produce an image from a seed (**random noise**)

```
def make_generator_model():
    model = tf.keras.Sequential()
    model.add(layers.Dense(7*7*256, use_bias=False, input_shape=(100,)))
    model.add(layers.BatchNormalization())
    model.add(layers.LeakyReLU())

    model.add(layers.Reshape((7, 7, 256)))
    assert model.output_shape == (None, 7, 7, 256) # Note: None is the batch size

    model.add(layers.Conv2DTranspose(128, (5, 5), strides=(1, 1), padding='same', use_bias=False))
    assert model.output_shape == (None, 7, 7, 128)
    model.add(layers.BatchNormalization())
    model.add(layers.LeakyReLU())

    model.add(layers.Conv2DTranspose(64, (5, 5), strides=(2, 2), padding='same', use_bias=False))
    assert model.output_shape == (None, 14, 14, 64)
    model.add(layers.BatchNormalization())
    model.add(layers.LeakyReLU())

    model.add(layers.Conv2DTranspose(1, (5, 5), strides=(2, 2), padding='same', use_bias=False, activation='tanh'))
    assert model.output_shape == (None, 28, 28, 1)

    return model
```

# Deep Convolutional GAN (DCGAN)

- Discriminator: a classifier

```
def make_discriminator_model():
    model = tf.keras.Sequential()
    model.add(layers.Conv2D(64, (5, 5), strides=(2, 2), padding='same',
                          input_shape=[28, 28, 1]))
    model.add(layers.LeakyReLU())
    model.add(layers.Dropout(0.3))

    model.add(layers.Conv2D(128, (5, 5), strides=(2, 2), padding='same'))
    model.add(layers.LeakyReLU())
    model.add(layers.Dropout(0.3))

    model.add(layers.Flatten())
    model.add(layers.Dense(1))

    return model
```

# Deep Convolutional GAN (DCGAN)

- Loss function: optimization goal
  - Discriminator loss: how well the discriminator can distinguish real images from fakes
  - Generator loss: how well it was able to trick the discriminator

```
def discriminator_loss(real_output, fake_output):  
    real_loss = cross_entropy(tf.ones_like(real_output), real_output)  
    fake_loss = cross_entropy(tf.zeros_like(fake_output), fake_output)  
    total_loss = real_loss + fake_loss  
    return total_loss  
  
def generator_loss(fake_output):  
    return cross_entropy(tf.ones_like(fake_output), fake_output)
```

# Deep Convolutional GAN (DCGAN)

Architecture guidelines for stable Deep Convolutional GANs

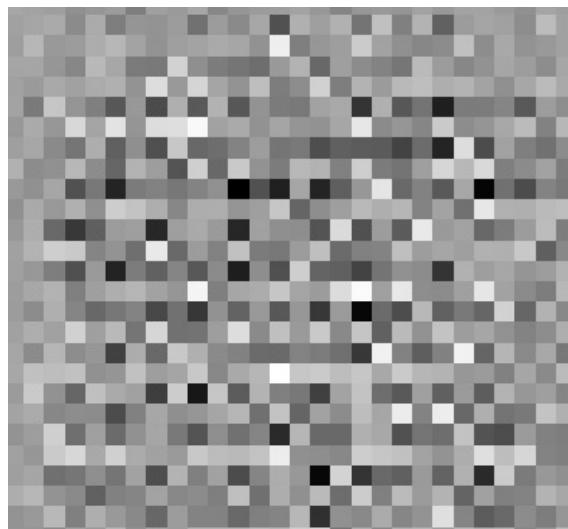
- Replace any pooling layers with strided convolutions (discriminator) and fractional-strided convolutions (generator).
- Use batchnorm in both the generator and the discriminator.
- Remove fully connected hidden layers for deeper architectures.
- Use ReLU activation in generator for all layers except for the output, which uses Tanh.
- Use LeakyReLU activation in the discriminator for all layers.

The GAN Zoo: <https://github.com/hindupuravinash/the-gan-zoo>

Tricks to make GAN better: <https://github.com/soumith/ganhacks>

Radford et al, “Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks”, ICLR 2016

# Effect of DCGAN



Start from: Random Noise



Synthesized Image

# Conditional GAN (cGAN)

- Learns a mapping from input images to output images
- cGAN: **Condition on input images** and generate corresponding output images

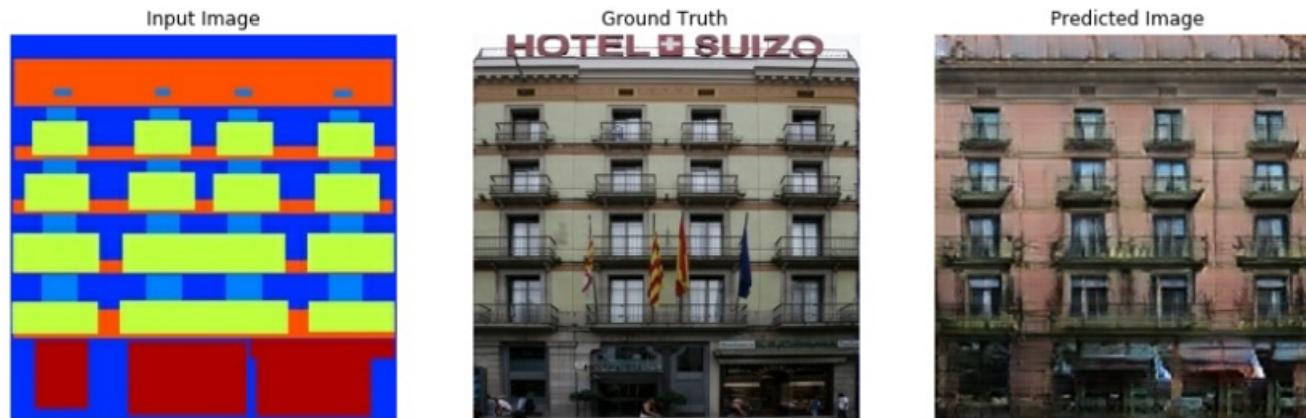
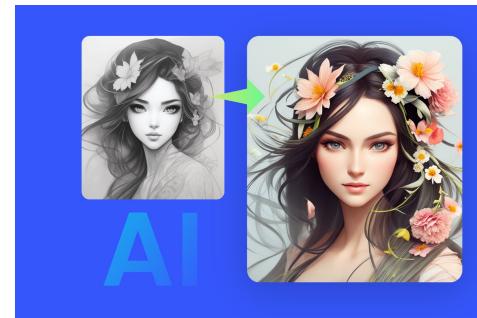


Image-to-Image Translation with Conditional Adversarial Networks (CVPR 2017)

# Applications of cGAN

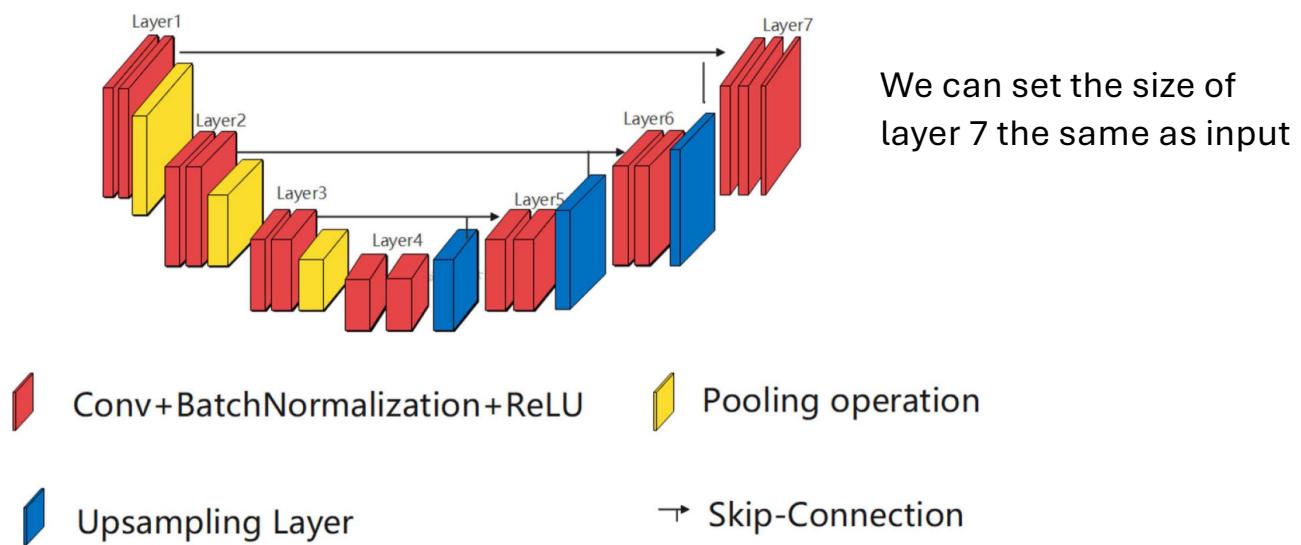
- Synthesizing photos from label maps
- Generating colorized photos from black and white images
- Turning Google Maps photos into aerial images
- Transforming sketches into photos...



From Dali museum

# Conditional GAN (cGAN)

- Generator (UNet): an encoder (downsampler) and decoder (upsampler)



<https://www.frontiersin.org/journals/aging-neuroscience/articles/10.3389/fnagi.2022.841297/full>

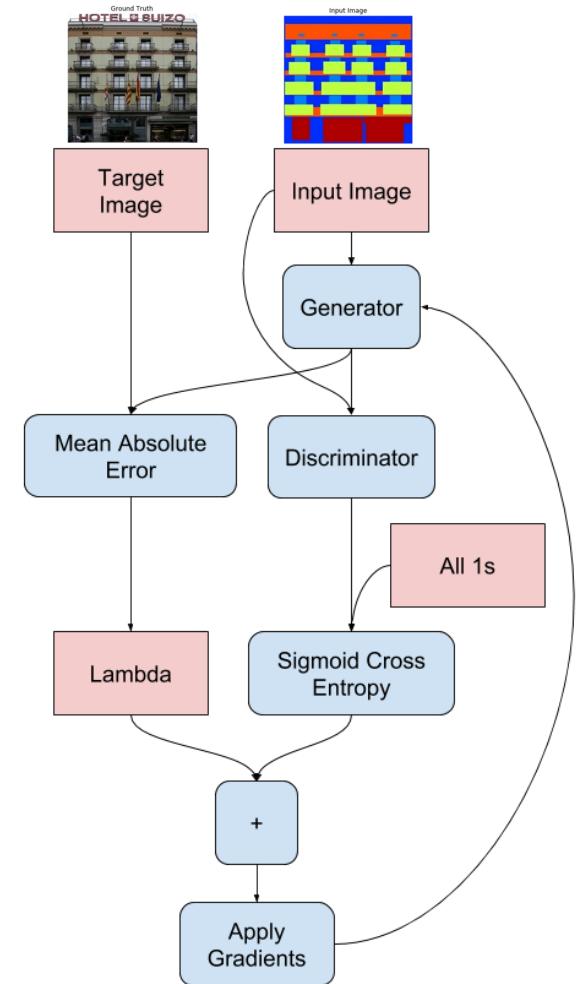
# Training of Generator

```
def generator_loss(disc_generated_output, gen_output, target):
    gan_loss = loss_object(tf.ones_like(disc_generated_output), disc_generated_output)

    # Mean absolute error
    l1_loss = tf.reduce_mean(tf.abs(target - gen_output))

    total_gen_loss = gan_loss + (LAMBDA * l1_loss)

    return total_gen_loss, gan_loss, l1_loss
```



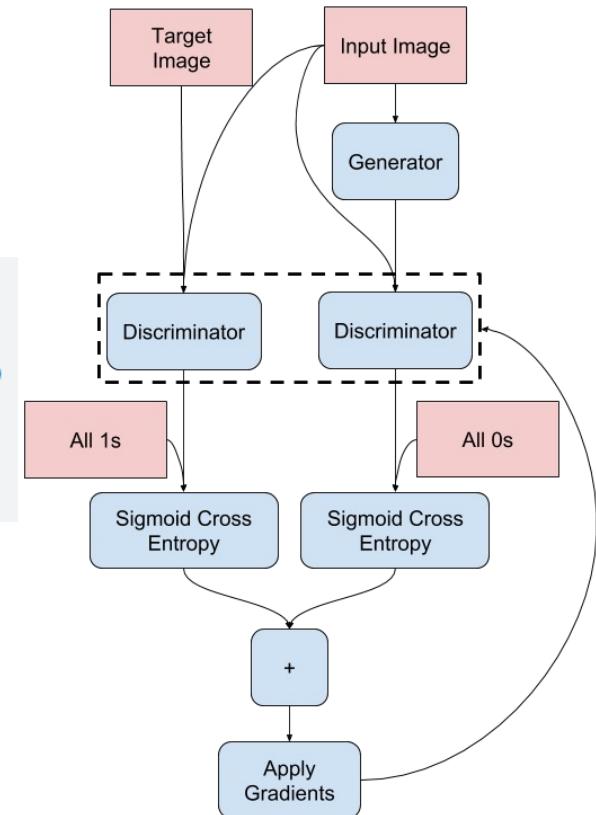
# Train of Discriminator

```
def discriminator_loss(disc_real_output, disc_generated_output):
    real_loss = loss_object(tf.ones_like(disc_real_output), disc_real_output)

    generated_loss = loss_object(tf.zeros_like(disc_generated_output), disc_generated_output)

    total_disc_loss = real_loss + generated_loss

    return total_disc_loss
```

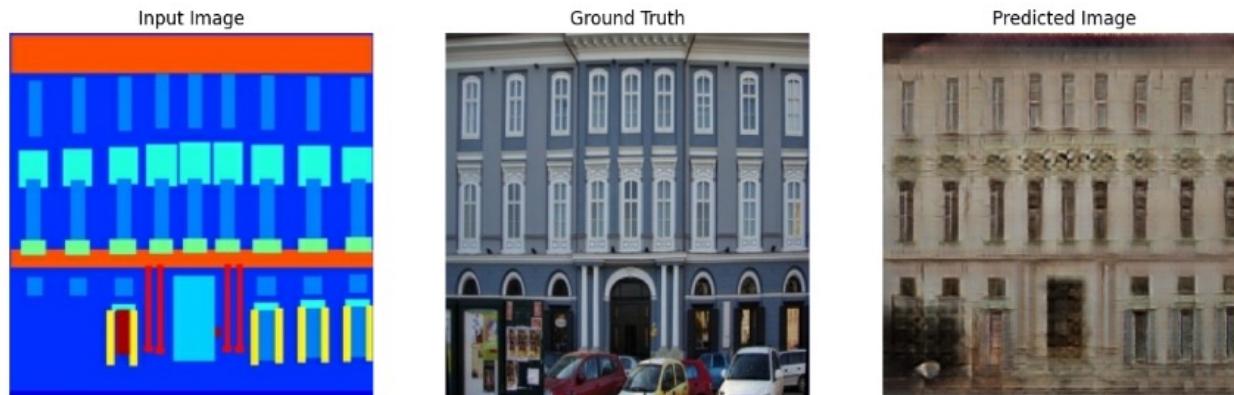


# Discriminator in cGAN

- Discriminator: a **convolutional PatchGAN** classifier—it tries to classify if each **image patch** is real or fake.
- The input image and the target image, which it should classify as real.
- The input image and the generated image (the output of the generator), which it should classify as fake.

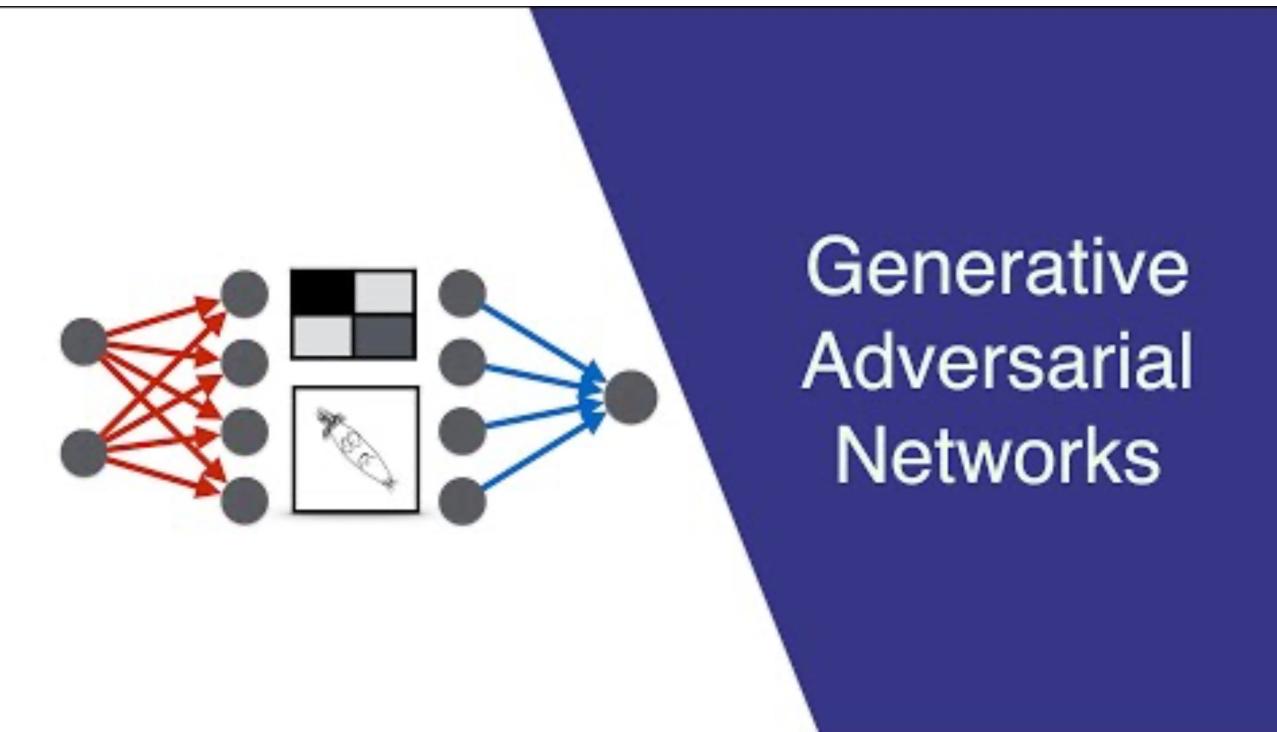
# Effect of cGAN (Pixel2Pixel)

- Pass images from the test set to the generator.
- The generator will then translate the input image into the output.



<https://www.tensorflow.org/tutorials/generative/pix2pix>

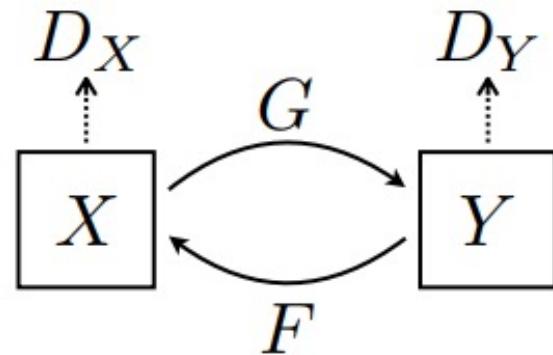
# Take a Break



<https://www.youtube.com/watch?v=8L11aMN5KY8>

# CycleGAN

There are 2 generators ( $G$  and  $F$ ) and 2 discriminators ( $X$  and  $Y$ ) being trained here.

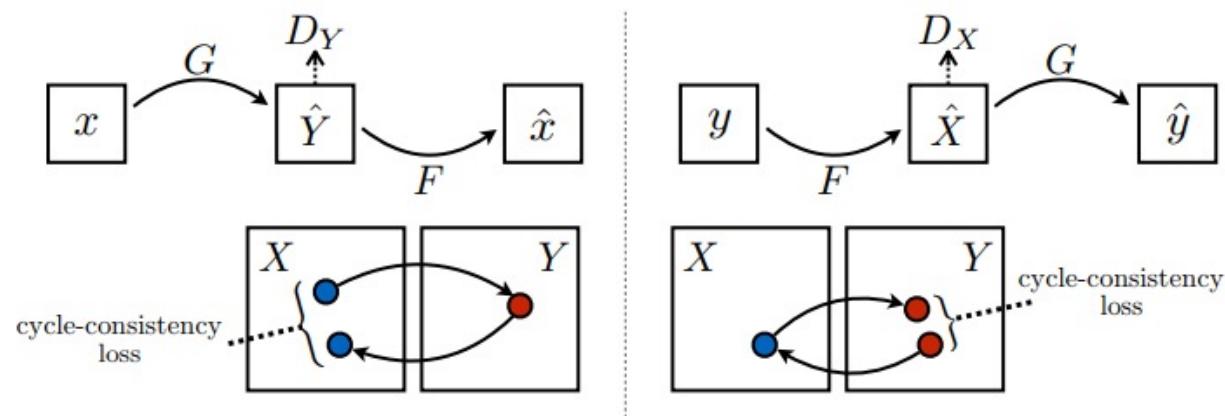


# CycleGAN

- Pixel2Pixel needs paired training data.
- CycleGAN: unpaired training data.
- CycleGAN uses instance normalization instead of batch normalization.
- The CycleGAN paper uses a modified Resnet-based Generator

# Loss Function in CycleGAN

- There is no pair data to train on, so cycle consistency loss is designed to enforce the network to learn meaningful mapping.
- **Cycle consistency** means the result should be close to the original input.



# Neural Style Transfer

- A *content image* and a *style reference image* (such as an artwork by a famous painter)
- Blend them together so the output image looks like the content image, but “painted” in the style of the style reference image.



Content Image



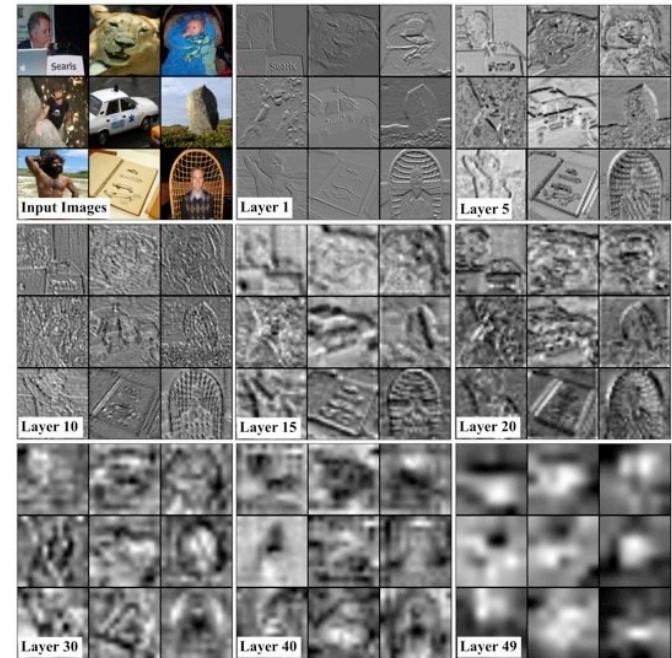
Style Image



Synthesized Image

# Content and Style Representations

- Use the **intermediate layers** of the model to get the *content* and *style* representations of the image.
- From edges, corners, textures to high-level concepts.



[https://www.researchgate.net/figure/Visualization-of-example-features-of-layers-1-10-20-30-40-and-49-of-a-deep\\_fig1\\_319622441](https://www.researchgate.net/figure/Visualization-of-example-features-of-layers-1-10-20-30-40-and-49-of-a-deep_fig1_319622441)

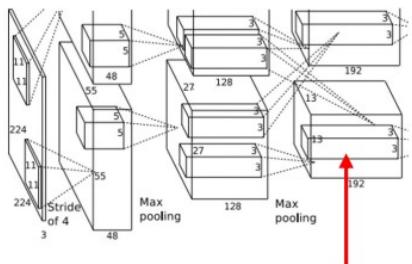
# Content and Style Representations

- The **content** of an image is represented by the values of the **intermediate feature maps**.
- The **style** of an image can be described by the **means and correlations** across the different feature maps.

# Style Representation: Gram Matrix



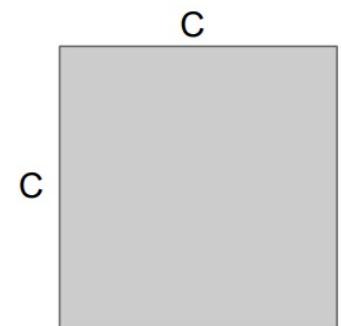
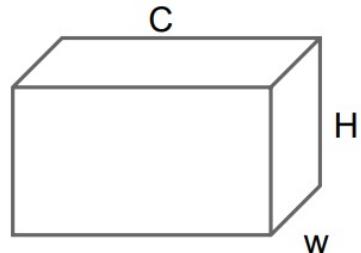
[This image](#) is in the public domain.



Each layer of CNN gives  $C \times H \times W$  tensor of features;  $H \times W$  grid of  $C$ -dimensional vectors

Outer product of two  $C$ -dimensional vectors gives  $C \times C$  matrix measuring co-occurrence

Average over all  $HW$  pairs of vectors, giving **Gram matrix** of shape  $C \times C$



Efficient to compute; reshape features from

$C \times H \times W$  to  $=C \times HW$

then compute  $G = FF^T$

Gram matrix: Ignore the positions of features and get correlations among features.

# Style Representation: Gram Matrix

The Gram matrix takes the outer product of the feature vector with itself at each location and averaging that outer product over all locations.

$$G_{cd}^l = \frac{\sum_{ij} F_{ijc}^l(x) F_{ijd}^l(x)}{IJ}$$

```
def gram_matrix(input_tensor):
    result = tf.linalg.einsum('bijc,bijd->bcd', input_tensor, input_tensor)
    input_shape = tf.shape(input_tensor)
    num_locations = tf.cast(input_shape[1]*input_shape[2], tf.float32)
    return result/(num_locations)
```

Tensorflow Implementation

```
def gram_matrix(input):
    a, b, c, d = input.size()  # a=batch size(=1)
    # b=number of feature maps
    # (c,d)=dimensions of a f. map (N=c*d)

    features = input.view(a * b, c * d)  # resize F_XL into |hat F_XL

    G = torch.mm(features, features.t())  # compute the gram product

    # we 'normalize' the values of the gram matrix
    # by dividing by the number of element in each feature maps.
    return G.div(a * b * c * d)
```

Pytorch Implementation

# Neural Style Transfer

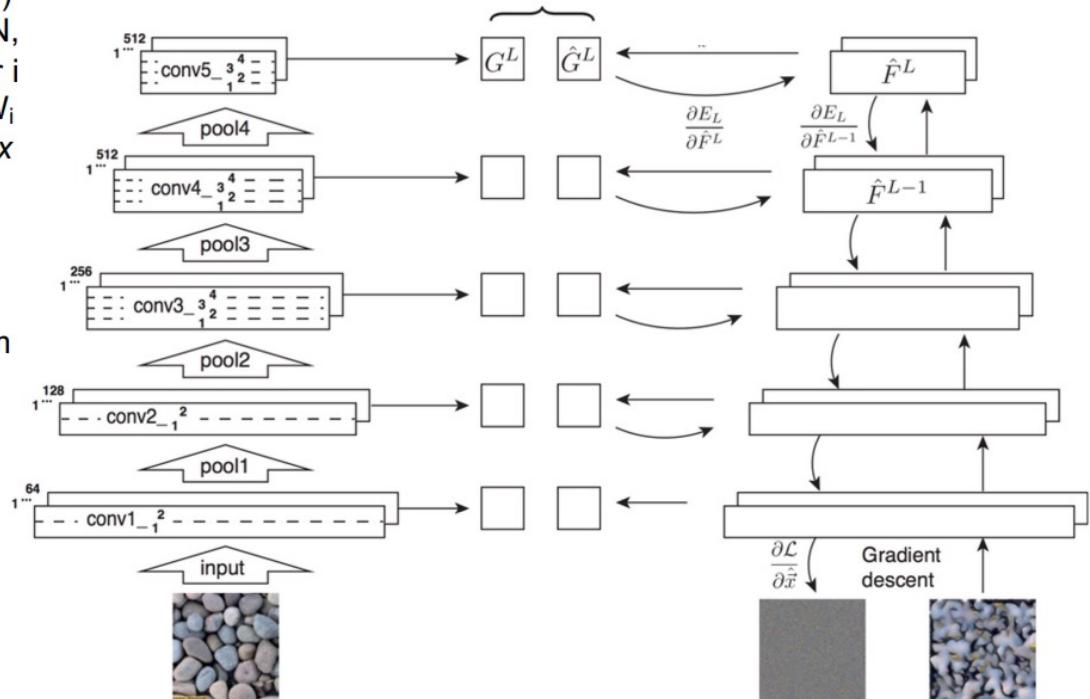
1. Pretrain a CNN on ImageNet (VGG-19)
2. Run input texture forward through CNN, record activations on every layer; layer i gives feature map of shape  $C_i \times H_i \times W_i$
3. At each layer compute the *Gram matrix* giving outer product of features:

$$G_{ij}^l = \sum_k F_{ik}^l F_{jk}^l \text{ (shape } C_i \times C_i\text{)}$$

4. Initialize generated image from random noise
5. Pass generated image through CNN, compute Gram matrix on each layer
6. Compute loss: weighted sum of L2 distance between Gram matrices
7. Backprop to get gradient on image
8. Make gradient step on image
9. GOTO 5

$$E_l = \frac{1}{4N_l^2 M_l^2} \sum_{i,j} (G_{ij}^l - \hat{G}_{ij}^l)^2$$

$$\mathcal{L}(\vec{x}, \hat{\vec{x}}) = \sum_{l=0}^L w_l E_l$$



# Learning Objective: MSE loss

```
def style_content_loss(outputs):
    style_outputs = outputs['style']
    content_outputs = outputs['content']
    style_loss = tf.add_n([tf.reduce_mean((style_outputs[name]-style_targets[name])**2)
                          for name in style_outputs.keys()])
    style_loss *= style_weight / num_style_layers

    content_loss = tf.add_n([tf.reduce_mean((content_outputs[name]-content_targets[name])**2)
                           for name in content_outputs.keys()])
    content_loss *= content_weight / num_content_layers
    loss = style_loss + content_loss
    return loss
```

```
@tf.function()
def train_step(image):
    with tf.GradientTape() as tape:
        outputs = extractor(image)
        loss = style_content_loss(outputs)

    grad = tape.gradient(loss, image)
    opt.apply_gradients([(grad, image)])
    image.assign(clip_0_1(image))
```

Gradient computation in TensorFlow

9/8/25

```
# We want to optimize the input and not the model parameters so we
# update all the requires_grad fields accordingly
input_img.requires_grad_(True)
# We also put the model in evaluation mode, so that specific layers
# such as dropout or batch normalization layers behave correctly.
model.eval()
model.requires_grad_(False)

optimizer = get_input_optimizer(input_img)
```

Gradient computation in PyTorch

CAI6605 Trustworthy AI Systems

35

# Neural Style Transfer



More weight to  
content loss

More weight to  
style loss

# References

- [https://cs231n.stanford.edu/slides/2024/lecture\\_11.pdf](https://cs231n.stanford.edu/slides/2024/lecture_11.pdf)
- [https://www.tensorflow.org/tutorials/generative/style\\_transfer](https://www.tensorflow.org/tutorials/generative/style_transfer)
- [https://pytorch.org/tutorials/advanced/neural\\_style\\_tutorial.html](https://pytorch.org/tutorials/advanced/neural_style_tutorial.html)
- <https://www.tensorflow.org/tutorials/generative/dcgan>
- <https://www.tensorflow.org/tutorials/generative/pix2pix>
- <https://www.tensorflow.org/tutorials/generative/cyclegan>