

# Trustworthy AI Systems

-- Large Language Model Agent

Instructor: Guangjing Wang  
[guangjingwang@usf.edu](mailto:guangjingwang@usf.edu)

# Last Lecture

- Recurrent Neural Network
- Attention
- Transformers
- Pretrained Foundation Model

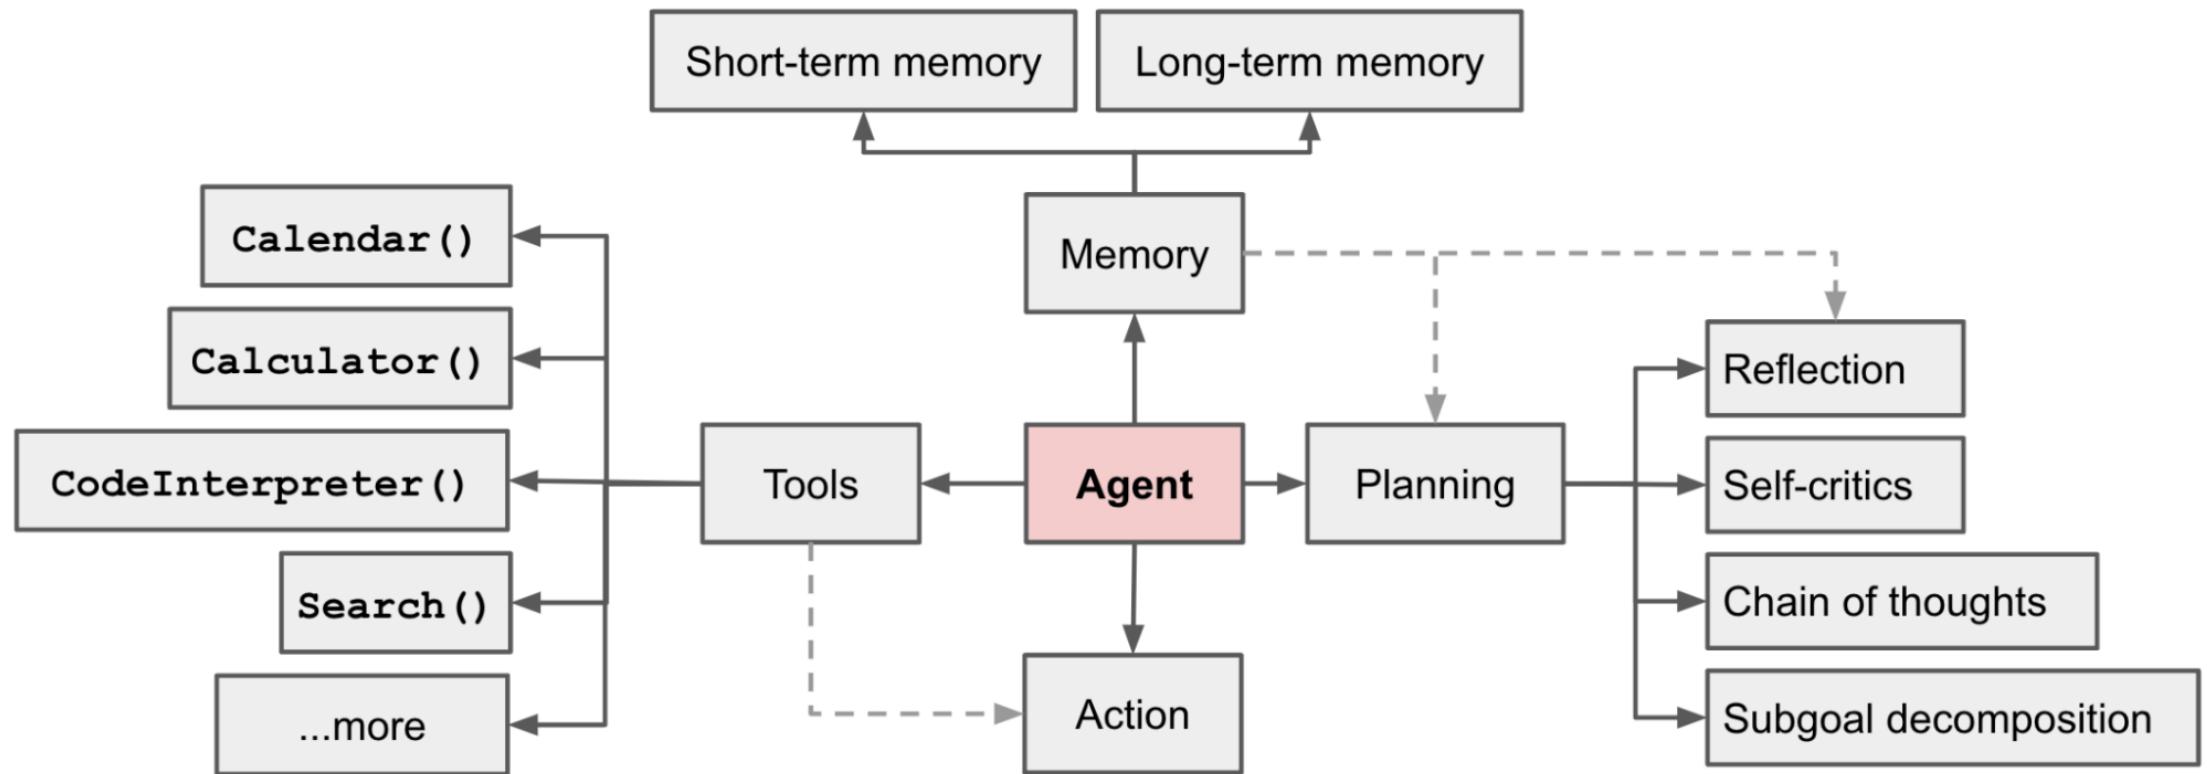
# This Lecture

- LLM Agent
- External Lecture: LLM Agents: Brief History and Overview

# LLM Agent

- A system that can use a Large Language Model (LLM) to reason through a problem, create a plan to solve the problem, and execute the plan with the help of a set of tools.
- LLM model example: <https://github.com/meta-llama/llama-models/tree/main>
- LLM Agent can respond to the inquiry by using planning, tailored focus, memory, using different tools, and breaking down a complex question into simpler sub-parts.

# LLM-powered Agent System



<https://lilianweng.github.io/posts/2023-06-23-agent/>

# An Application Example

- Consider a LLM application that is designed to help financial analysts answer questions about the performance of a company.
  - “What were the three takeaways from the Q2 earnings call from FY23? Focus on the technological moats that the company is building”.

How do we develop a solution to answer a question like above?

# Agent Core

- The agent core is the central coordination module that manages the core logic and behavioral characteristics of an Agent.

```
template = """GENERAL INSTRUCTIONS
Your task is to answer questions. If you cannot answer the question, request a
helper or use a tool. Fill with Nil where no tool or helper is required.

AVAILABLE TOOLS
- Search Tool
- Math Tool

AVAILABLE HELPERS
- Decomposition: Breaks Complex Questions down into simpler subparts

CONTEXTUAL INFORMATION
<No previous questions asked>

QUESTION
How much did the revenue grow between Q1 of 2024 and Q2 of 2024?

ANSWER FORMAT
{"Tool_Request": "<Fill>", "Helper_Request": "<Fill>"}""""
```

Basic template of how the different modules of an agent are assembled in its core.

# Planning Module

- Task and question decomposition
  - “What were the three takeaways from NVIDIA’s last earnings call?”
    - “Which technological shifts were discussed the most?”
    - “Are there any business headwinds?”
    - “What were the financial results?”
  - Each of these questions can be further broken into subparts.
  - A specialized AI agent must guide this decomposition.
- Self-Reflection: Techniques such as Chain of Thought, and Graph of thought have served as critic– or evidence-based prompting frameworks.

# Chain-of-Thought (CoT)

- A series of intermediate reasoning steps to significantly improve the ability of LLM to perform complex reasoning.

(a) Few-shot

Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?

A: The answer is 11.

Q: A juggler can juggle 16 balls. Half of the balls are golf balls, and half of the golf balls are blue. How many blue golf balls are there?

A:

(Output) The answer is 8. ✗

(b) Few-shot-CoT

Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?

A: Roger started with 5 balls. 2 cans of 3 tennis balls each is 6 tennis balls.  $5 + 6 = 11$ . The answer is 11.

Q: A juggler can juggle 16 balls. Half of the balls are golf balls, and half of the golf balls are blue. How many blue golf balls are there?

A:

(Output) The juggler can juggle 16 balls. Half of the balls are golf balls. So there are  $16 / 2 = 8$  golf balls. Half of the golf balls are blue. So there are  $8 / 2 = 4$  blue golf balls. The answer is 4. ✓

(c) Zero-shot

Q: A juggler can juggle 16 balls. Half of the balls are golf balls, and half of the golf balls are blue. How many blue golf balls are there?

A: The answer (arabic numerals) is

(Output) 8 ✗

(d) Zero-shot-CoT (Ours)

Q: A juggler can juggle 16 balls. Half of the balls are golf balls, and half of the golf balls are blue. How many blue golf balls are there?

A: Let's think step by step.

(Output) There are 16 balls in total. Half of the balls are golf balls. That means that there are 8 golf balls. Half of the golf balls are blue. That means that there are 4 blue golf balls. ✓

# Memory Module

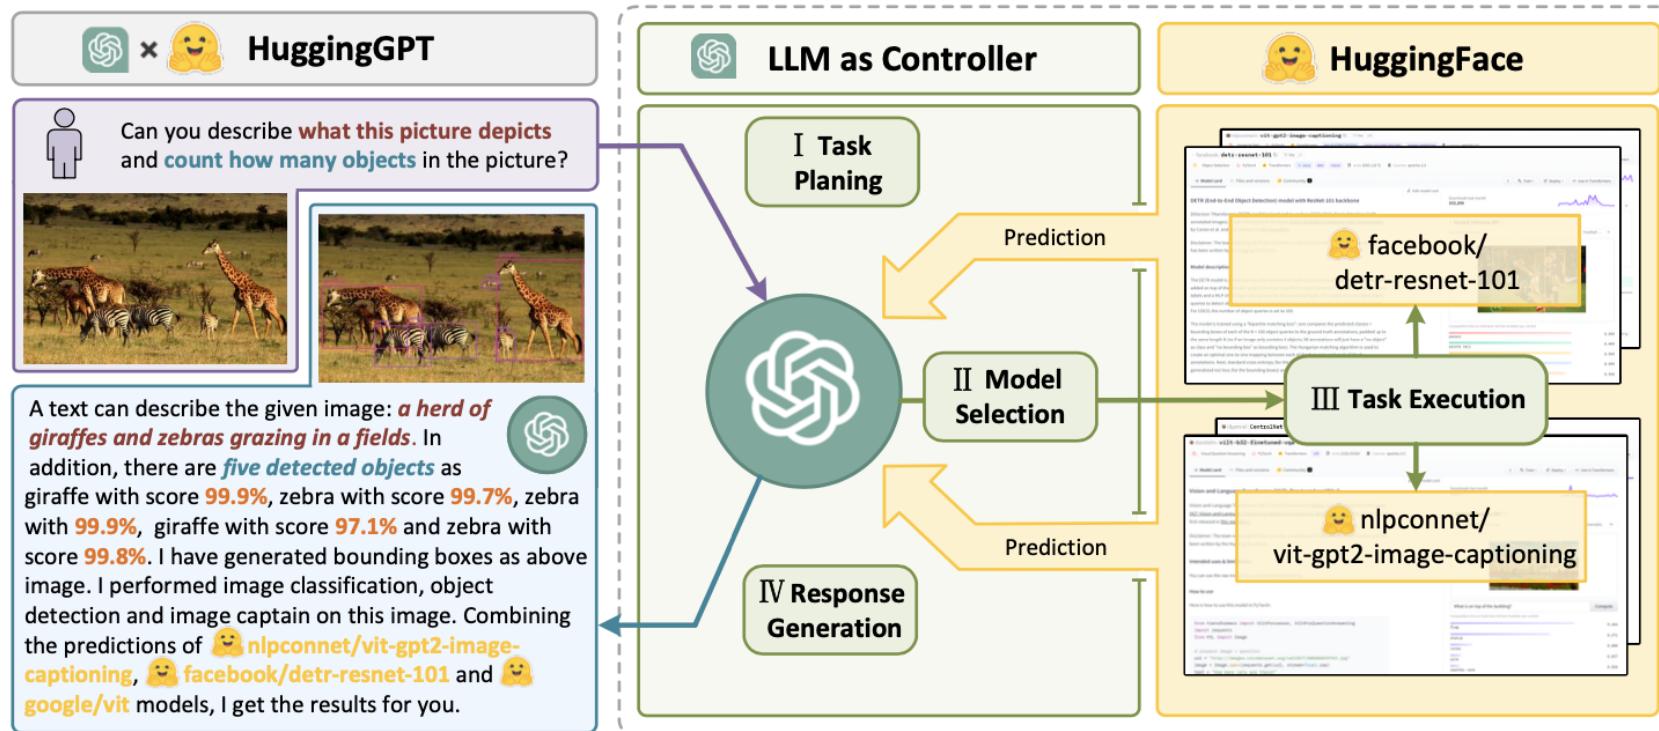
- **Short-term memory:** A ledger of actions and thoughts that an agent goes through to attempt to answer a single question from a user: the agent’s “train of thought.” (e.g., in-context learning)
- **Long-term memory:** A ledger of actions and thoughts about events that happen between the user and agent. It is a log book that contains a conversation history stretching across weeks or months.

Memory requires a composite score, which is made up of semantic similarity, importance, recency, and other application-specific metrics.

# Tools

- Tools are well-defined executable workflows that agents can use to execute tasks
  - A Retrieval Augmented Generation (RAG) pipeline to generate context aware answers
  - A code interpreter to solve complex programmatically tasks
  - An API to search information over the internet
  - Any simple API service like a weather API or an API for an Instant messaging application
  - ...

# Example: HuggingGPT



Using ChatGPT as the task planner to select models available in HuggingFace platform according to the model descriptions and summarize the response based on the execution results.

# HuggingGPT: Task Planning

- Using few-shot examples to guide LLM to do task parsing and planning.

The AI assistant can parse user input to several tasks: `[{"task": task, "id": task_id, "dep": dependency_task_ids, "args": {"text": text, "image": URL, "audio": URL, "video": URL}}]`. The "dep" field denotes the id of the previous task which generates a new resource that the current task relies on. A special tag "-task\_id" refers to the generated text image, audio and video in the dependency task with id as task\_id. The task MUST be selected from the following options: {{ Available Task List }}. There is a logical relationship between tasks, please note their order. If the user input can't be parsed, you need to reply empty JSON. Here are several cases for your reference: {{ Demonstrations }}. The chat history is recorded as {{ Chat History }}. From this chat history, you can find the path of the user-mentioned resources for your task planning.

# HuggingGPT: Model Selection

- LLM distributes the tasks to expert models.

Given the user request and the call command, the AI assistant helps the user to select a suitable model from a list of models to process the user request. The AI assistant merely outputs the model id of the most appropriate model. The output must be in a strict JSON format: "id": "id", "reason": "your detail reason for the choice". We have a list of models for you to choose from {{ Candidate Models }}. Please select one model from the list.

# HuggingGPT: Task Execution

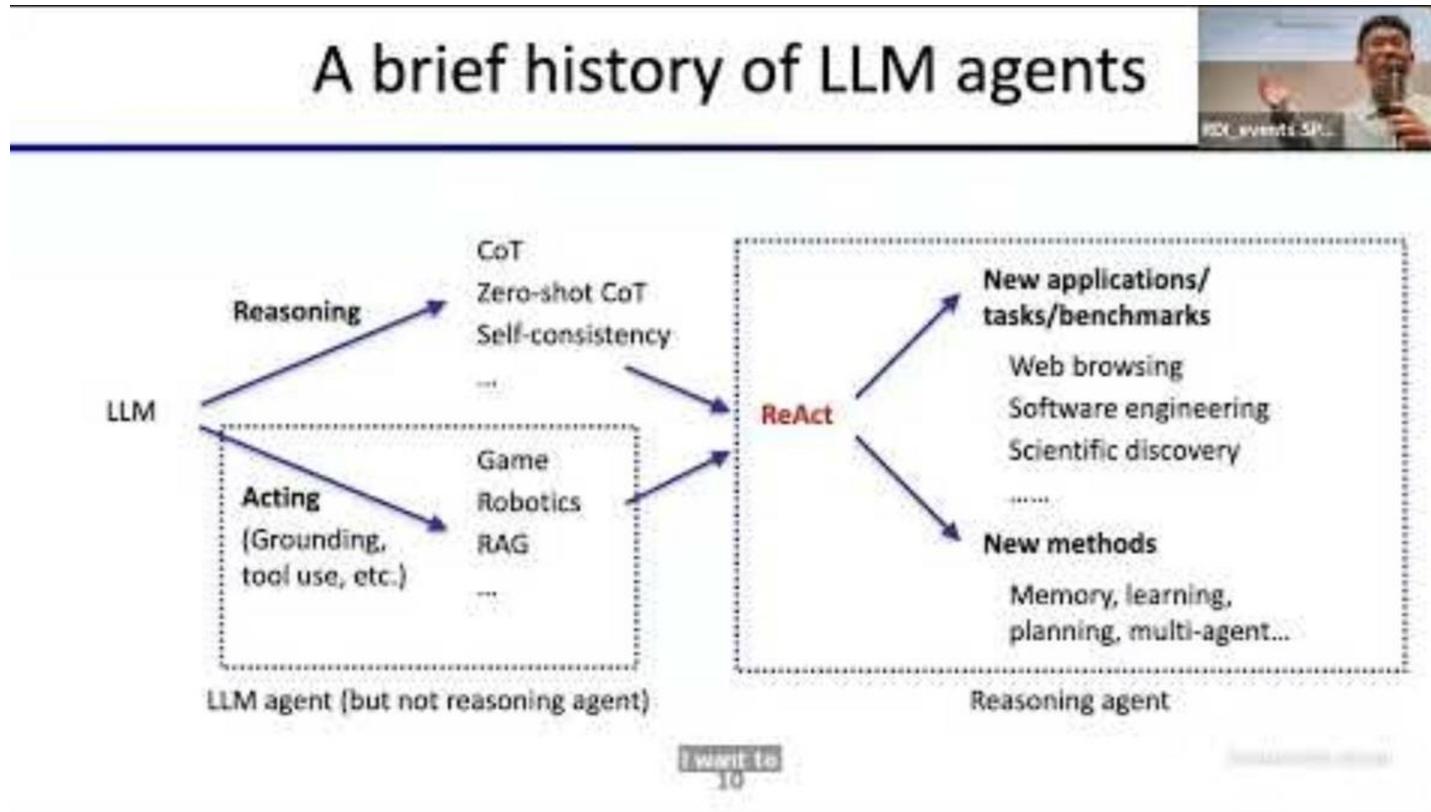
- Expert models execute on the specific tasks and log results.

With the input and the inference results, the AI assistant needs to describe the process and results. The previous stages can be formed as - User Input: {{ User Input }}, Task Planning: {{ Tasks }}, Model Selection: {{ Model Assignment }}, Task Execution: {{ Predictions }}. You must first answer the user's request in a straightforward manner. Then describe the task process and show your analysis and model inference results to the user in the first person. If inference results contain a file path, must tell the user the complete file path.

# Challenges for Practical Usage

- Efficiency improvement is needed as both LLM inference rounds and interactions with other models slow down the process;
- It relies on a long context window to communicate over complicated task content;
- Stability improvement of LLM outputs and external model services.
- The reliability of model outputs is questionable, as LLMs may make formatting errors and occasionally exhibit rebellious behavior (e.g. refuse to follow an instruction)

# LLM Agents: Brief History and Overview



<https://www.youtube.com/watch?v=RM6ZArd2nVc>

# References

- <https://developer.nvidia.com/blog/introduction-to-lm-agents/>
- <https://www.promptingguide.ai/>
- <https://lilianweng.github.io/posts/2023-06-23-agent/>
- <https://nips.cc/virtual/2023/tutorial/73948>