

Trustworthy AI Systems

-- Privacy of AI

Instructor: Guangjing Wang
guangjingwang@usf.edu

Last Lecture


- Poisoning Attacks
- Poisoning Scenarios
 - Centralized
 - Distributed
- Defense for Poisoning Attacks

This Lecture

- Membership Inference Attacks
- Model Inversion Attacks
- Model Stealing Attacks
- Privacy Protection Methods

Membership Inference Attacks

- Determine whether an individual data instance x^* is part of the training dataset \mathcal{D} for a model.
- The membership inference attacks on both **supervised** classification models and **generative** models (GANs, VAEs) have been demonstrated.
- A common approach is to first train several **shadow models** that **imitate** the behavior of the target model and use the **prediction vectors** of the shadow models to train a binary classifier (that infers the membership).



<https://dl.acm.org/doi/abs/10.1145/3436755>


Shadow Training Attack (1)

- Threat model:
 - The adversary has **back-box** query access to the target model
 - The goal is to infer whether input samples were part of its private training set
- Shadow training approach:
 - Create several shadow models to substitute the target model
 - Each shadow model is trained on a dataset that has a similar distribution as the private training dataset of the target model
 - E.g., if the target model performs celebrity face recognition, the attacker can collect images of celebrities from the Internet

Shadow Training Attack (2)

- The output probability vectors from the shadow models are next used as inputs to train attack models (as binary classifiers) for each class
 - E.g., the probability vectors for all input images of Alice from **all shadow training sets** are labeled with 1 (meaning ‘in’ the training set)
 - The probability vectors for all input images of Alice from **all shadow test sets** are labeled with 0 (meaning ‘out’ or not in the training set)
 - An attack model is trained on these inputs to perform binary classification (in or out)
 - A separate attack model is trained for each celebrity person in the shadow training sets

Shadow Training Attack (3)



<https://arxiv.org/abs/1610.05820>

Shadow Training Attack (4)

- The attack models for each class are afterward used to predict whether individual input instances were members of the private training set of the target model.
- The assumption in this attack is that the output probability vectors for samples that are members of the training sets are different from samples out of the training sets.
- Experiments showed that increasing the number of shadow models improves the accuracy of membership inference, but it also increases the computational resources.

This Lecture

- Membership Inference Attacks
- Model Inversion Attacks
- Model Stealing Attacks
- Privacy Protection Methods

Model Inversion Attack (1)

- *Model inversion attack* creates prototype examples for the classes in the dataset
 - The authors demonstrated an attack against a DNN model for face recognition.
 - Given a **person's name (label)** and **white-box access to the model**, the attack reverse-engineered the model and produced an averaged image of that person.
 - The obtained averaged image (left image below) makes the person recognizable.
 - This attack is limited to classification models where the classes pertain to one type of object (such as the faces of the same person).

Recovered Image



Training Image Data



<https://dl.acm.org/doi/10.1145/2810103.2813677>

Model Inversion Attack (2)

- The model inversion attack applies gradient descent to start from a given label and follows the gradient in a trained network to recreate an image for that label
 - In the algorithm, c denotes the cost function, whereas the PROCESS function applies image denoising and sharpening operations to improve the reconstructed image

Algorithm 1 Inversion attack for facial recognition models.

```
1: function MI-FACE(label,  $\alpha$ ,  $\beta$ ,  $\gamma$ ,  $\lambda$ )
2:    $c(\mathbf{x}) \stackrel{\text{def}}{=} 1 - \tilde{f}_{\text{label}}(\mathbf{x}) + \text{AUXTERM}(\mathbf{x})$ 
3:    $\mathbf{x}_0 \leftarrow 0$ 
4:   for  $i \leftarrow 1 \dots \alpha$  do
5:      $\mathbf{x}_i \leftarrow \text{PROCESS}(\mathbf{x}_{i-1} - \lambda \cdot \nabla c(\mathbf{x}_{i-1}))$ 
6:     if  $c(\mathbf{x}_i) \geq \max(c(\mathbf{x}_{i-1}), \dots, c(\mathbf{x}_{i-\beta}))$  then
7:       break
8:     if  $c(\mathbf{x}_i) \leq \gamma$  then
9:       break
10:    return  $[\arg \min_{\mathbf{x}_i} (c(\mathbf{x}_i)), \min_{\mathbf{x}_i} (c(\mathbf{x}_i))]$ 
```

AuxTerm: case-specific function, any available auxiliary information to inform the cost function.

$f_{\{\text{label}\}}$: facial recognition model

GAN-based Model Inversion Attack

Inferring sensitive features (e.g., face) in the training data:

Rather than reconstructing private training data from scratch, we leverage partial public information, to **learn a distributional prior** via generative adversarial networks (GANs) and use it to guide the inversion process.




Figure 1: Overview of the proposed GMI attack method.

GAN-based Model Inversion Attack

Stage 1: Train the generator and the discriminators on public datasets in order to encourage the generator to generate realistic-looking images. F is the feature extractor of the Target Network, L_{div} is to promotes the diversity of generated images

$$L_{\text{wgan}}(G, D) = E_x[D(x)] - E_z[D(G(z))]$$

$$L_{\text{div}}(G) = E_{\mathbf{z}_1, \mathbf{z}_2} \left[\frac{\|F(G(\mathbf{z}_1)) - F(G(\mathbf{z}_2))\|}{\|\mathbf{z}_1 - \mathbf{z}_2\|} \right]$$

Stage 2: Find the latent vector that generates an image achieving the maximum likelihood under the target network while remaining realistic.

$$\hat{z} = \arg \min_z L_{\text{prior}}(z) + \lambda_i L_{\text{id}}(z)$$

$$L_{\text{prior}}(z) = -D(G(z)) \quad L_{\text{id}}(z) = -\log[C(G(z))]$$




Figure 1: Overview of the proposed GMI attack method.

This Lecture

- Membership Inference Attacks
- Model Inversion Attacks
- Model Stealing Attacks
- Privacy Protection Methods

Model Stealing Attack (Model Extraction Attack)

- Adversarial goal: reconstruct an approximated model $f'(x)$ of the target model $f(x)$.
- The approximated function $f'(x)$ will act as a substitute model and produce similar outputs as the target model.
 - The adversary has black-box query access to the model
 - The goal is to “steal” the model and use the substitute model for launching other attacks, such as synthesis of adversarial examples, or membership inference attacks
- Besides creating a substitute model, several works focused on recovering the hyperparameters of the model, such as the number of layers, optimization algorithm, activation types used, etc.

This Lecture

- Membership Inference Attacks
- Model Inversion Attacks
- Model Stealing Attacks
- Privacy Protection Methods

Causes of Privacy Leaks in Machine Learning

- Overfitting
 - It leads to poor generalization and memorization of the training data
 - Although **adversarial training** is often applied for increasing model robustness, it reduces the accuracy of model on clean data, due to the trade-off between the model accuracy and robustness
 - The reduced accuracy can lead to increased sensitivity to data leakage
- Datasets that are **more diverse** and with a larger number of categories are **more susceptible** to attacks
 - Binary classifiers are safer than multiclass models
 - Input samples that are out-of-distribution (i.e., are considered **outliers** with respect to the distribution of the training data) are more susceptible to privacy leakage
- Model complexity
 - Complex models with a large number of parameters memorize more sensitive information about the training data

Defenses against Privacy Attacks

- Anonymization techniques
- Encryption techniques
- Differential privacy
- Distributed learning
- ML-specific techniques

Data Privacy Protection

- Data privacy protection techniques have the goal of allowing analysts to learn about trends in data, without revealing information specific to individual data instances
 - Therefore, privacy techniques involve an intentional release of information, and an attempt to control what can be learned from the released information
- The Fundamental Law of Information Recovery states that “overly accurate estimates of too many statistics can completely destroy privacy”
 - I.e., extracting useful information from a dataset (e.g., for training an ML model) poses a privacy risk to the data
- There is an inevitable trade-off between privacy and accuracy (i.e., utility)
 - Preferred privacy techniques should provide an estimate of how much privacy is lost by interacting with data

Anonymization Techniques

- ***Anonymization*** techniques provide privacy protection by removing identifying information from the data
- E.g., remove Personally Identifiable Information (PII)
 - In the example below, the Name and Address columns are masked

User ID	Name	Address	Account Type	Subscription Date
001	Alice	123 A St	Pro	01/02/20
002	Bob	234 B St	Free	02/03/21
003	Charlie	456 C St	Pro	03/04/18

User ID	Name	Address	Account Type	Subscription Date
001	A**	*** A St	Pro	01/02/20
002	B**	*** B St	Free	02/03/21
003	C**	*** C St	Pro	03/04/18

Anonymization Techniques


- Drawback: The remaining information in the data can be used for identifying the individual data instances
 - For example, based on health records (including diagnoses and prescriptions) with removed personal information released by an insurance group in 1997, a researcher extracted the information for the Governor of Massachusetts
 - This is referred to as *de-anonymization*
 - The same researcher later showed that 87% of all Americans can be uniquely identified using 3 groups of information: ZIP code, birth date, and gender

K-anonymity

- *k-anonymity* is an approach for protecting data privacy by suppressing certain identifying data features
 - This approach removes fields of data for individuals who have unique characteristics
 - E.g., students at UI who are from Latvia and are enrolled in Architecture
- A dataset is *k-anonymous* if, for any person's record, there are at least $k - 1$ other records that are indistinguishable
 - Therefore, a linkage attack will result in a group of k records that can belong to a person of interest
- Limitation: this approach is mostly applicable to large datasets with low-dimensional input features
 - The more input features for each record, the higher the possibility of unique records

Encryption Techniques

- Encryption is a cryptography approach, which converts the original representation of information (plaintext) into an alternative form (Ciphertext)
 - The sender of encrypted information shares the decoding technique only with the intended recipients of the information



<https://www.twilio.com/en-us/blog/what-is-public-key-cryptography>

Encryption Techniques


- Encrypting the training data has been applied in ML
 - Common techniques for data encryption include:
 - Homomorphic encryption (HE)
 - Secure multi-party computation (SMPC)
- Encrypting ML models is a less common approach
 - Homomorphic encryption has been applied to the model gradients in a collaborative deep learning setting to protect the model privacy

Homomorphic Encryption

- Homomorphic encryption (HE) allows users to perform computations on encrypted data (without decrypting it)
 - Encrypted data can be analyzed and manipulated without revealing the original data
- HE uses a public key to encrypt the data and applies an algebraic system (e.g., additions and multiplications) to allow computations while the data is still encrypted
 - Only the person who has a matching private key can access the decrypted results


Homomorphic Encryption

- In ML, training data can be encrypted and sent to a server for model training.
 - Even if the server is untrusted or compromised, the confidentiality of the data will remain preserved.
 - One main limitation of HE is the slowing down of the training process.
- HE has been applied to traditional ML approaches.
 - Training DNNs over encrypted data is still challenging, due to the increased computational complexity.



Secure Multi-Party Computation

- Secure Multi-Party Computation (SMPC) is an extension of encryption in multi-party setting.
 - SMPC allows two or more parties to jointly perform computation over their private data, without sharing the data.
 - E.g., two banks want to know if they have both flagged the same individuals and learn about the activities of those individuals.
 - The banks can share encrypted tables of flagged individuals, and they can decrypt only the matched records, but not the information for individuals that are not in both tables.



<https://arxiv.org/abs/1909.11701>

Secure Multi-Party Computation

- In ML, SMPC can be used to compute updates of the **model parameters** by multiple parties **without** sharing their private data
 - For example, SMPC has been applied to federated learning, where participants **encrypt** their updates, and the central server can recover only the **sum** of the updates from all participants
 - Besides data privacy, SMPC also offers protection against adversarial participants
 - Either all parties are honest and can jointly compute the correct output, or if a malicious party is dishonest the joint output will be incorrect
- SMPC has been applied to traditional ML models, such as decision trees, linear regression, logistic regression, Naïve Bayes, k -means clustering
 - Application of SMPC to DNNs is also challenging, due to increased **computational costs**

SMPC and HE

- SMPC protects the **privacy** of the data in collaborative learning
 - E.g., participants in collaborative learning do not trust the other participants or the central server
- HE protects the **confidentiality** of the data from external adversaries
 - E.g., a data owner wants to use an **MLaaS (Machine Learning as a Service)**, but does not trust the service provider: the owner sends encrypted data, the provider processes encrypted data and sends back encrypted results, the owner decrypts the results
 - Or, a bank can store encrypted banking information in the cloud, and use HE to ensure that only the employees of the bank can access the data

Differential Privacy


- **Differential privacy** is based on employing obfuscation mechanisms for privacy protection
 - A **randomization mechanism $\mathcal{M}(D)$** applies noise ξ to the outputs of a function $f(D)$ to protect the privacy of individual data instances, i.e., $\mathcal{M}(D) = f(D) + \xi$
 - Commonly used randomization mechanisms include Laplacian, Gaussian, and Exponential mechanism
- DP is often implemented in practical applications, and examples include:
 - 2014: Google's RAPPOR, for statistics on unwanted software hijacking users' settings
 - 2015: Google, for sharing historical traffic statistics
 - 2016: Apple, for improving its Intelligent personal assistant technology
 - 2017: Microsoft, for telemetry in Windows
 - 2020: LinkedIn, for advertiser queries
 - 2020: U.S. Census Bureau, for demographic data

Differential Privacy

- In ML, DP is achieved by adding noise to:
 - *Model parameters*
 - Several works applied DP to conventional ML methods.
 - Differentially private SGD (Abadi, 2016) clips and adds noise to the gradients of deep NNs during training.
 - This reduces the memorization of individual input instances by the model.
 - The approaches that apply obfuscation to the model parameters via DP are also referred to as differentially private ML.
 - *Model outputs*
 - PATE (Private Aggregation of Teacher Ensembles) approach (Papernot, 2018) employs an ensemble of models trained on disjoint subsets of the training data, called teacher models.
 - Noise is added to the outputs of the teacher models, and the aggregated outputs are used to train another model, called the student model.
 - *Training data*
 - Obfuscation of training data in ML has been also investigated in several works.

Distributed Learning

- **Distributed learning** allows multiple parties to train a global model without releasing their private data
- Some form of **aggregation** is applied to the local updates of the model parameters by the users in distributed learning to create a global model
 - E.g., averaging is one common form of aggregation
- **Federated learning** is the most popular distributed learning scheme



<https://arxiv.org/abs/2011.11819>

Distributed Learning

- **Federated learning** or collaborative learning – learn one global model using data stored at multiple locations (e.g., remote devices)
 - The data are processed locally and used to update the model
 - The data does not leave the remote devices and remains private
 - The central server aggregates the updates and creates the global model
- Decentralized Peer-to-Peer (P2P) learning – the remote devices communicate and exchange the updates directly, without a central server
 - Removes the need to send updates to a potentially untrusted central server
- **Split learning** – each remote device is used to train several layers of the global model, and send the outputs to a central server
 - The remote devices can train the first several layers of a DNN, and the central server can train the remaining final layers
 - The gradient is back-propagated from the central server to each user to sequentially complete the back-propagation through all layers of the model
 - The devices send the outputs of intermediate layers, rather than model parameters
 - Split learning is more common for IoT devices with limited computational resources

ML-Specific Techniques

- Overfitting is one of the reasons for information leakage:
 - Regularization techniques in ML can therefore be used to reduce overfitting, as well as a defense strategy
 - Different regularization techniques in NNs include:
 - Explicit regularization: dropout, early stopping, weight decay;
 - Implicit regularization: batch normalization.
- Other ML-specific techniques include:
 - Dimensionality reduction – removing inputs with features that occur rarely in the training set;
 - Weight-normalization – rescaling the weights of the model during training;
 - Selective gradient sharing – in federated learning, the users share a fraction of the gradient at each update;
 - Data Reconstruction – Remove sensitive features in the data and keep the main utility-related features.

References

- When Machine Learning Meets Privacy: A Survey and Outlook (<https://dl.acm.org/doi/abs/10.1145/3436755>)
- Beyond Boundaries: A Comprehensive Survey of Transferable Attacks on AI Systems (<https://arxiv.org/abs/2311.11796>)
- A Survey on Federated Learning Systems: Vision, Hype and Reality for Data Privacy and Protection (<https://ieeexplore.ieee.org/abstract/document/9599369>)
- A Survey on Differential Privacy for Unstructured Data Content (<https://dl.acm.org/doi/full/10.1145/3490237>)
- <https://differentialprivacy.org>